

Introduction

Before going to the AWS console and creating an Auto Scaling Group, it's important to understand the key components of an Auto Scaling Group. AWS has done an excellent job defining them so the official definition is placed below for convenience sake:

Groups

Your EC2 instances are organized into groups so that they can be treated as a logical unit for the purposes of scaling and management. When you create a group, you can specify its minimum, maximum, and desired number of EC2 instances.

Launch configurations

Your group uses a launch configuration as a template for its EC2 instances. When you create a launch configuration, you can specify information such as the AMI ID, instance type, key pair, security groups, and block device mapping for your instances.

Launch template

A launch template is similar to a launch configuration, in that it specifies instance configuration information. ... However, defining a launch template instead of a launch configuration allows you to have multiple versions of a template. With versioning, you can create a subset of the full set of parameters and then reuse it to create other templates or template versions.

After you complete this Lab you can read the [full AWS documentation on Auto Scaling here](#).

Auto Scaling groups are commonly paired with load balancers that evenly distribute requests across all the instances in the group. In this Lab, you will create a load balancer before creating an Auto Scaling group so that the Auto Scaling group can be configured to work with the load balancer at creation time.